

Database System Architecture and Implementation

Module 0
Introduction and Overview
October 22, 2018

Module Overview

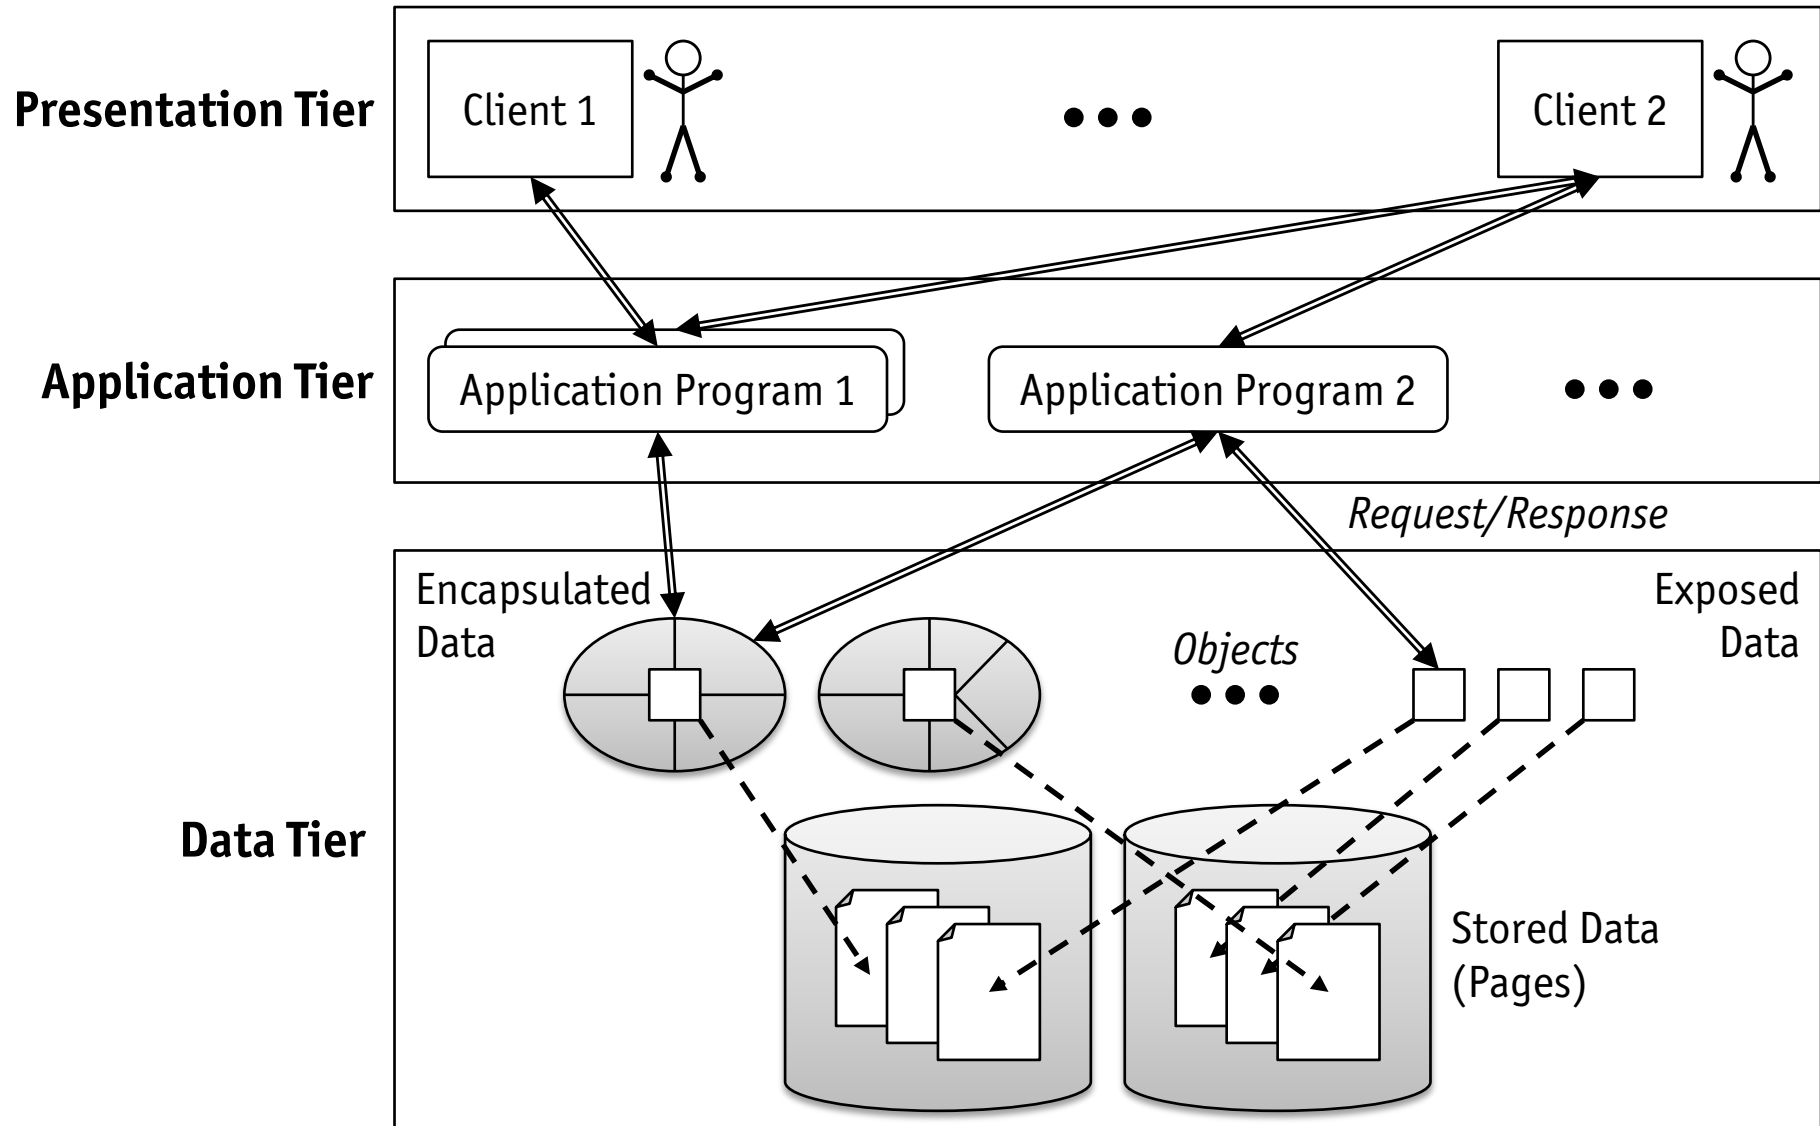
1. Course contents

1. Goals of this course
2. DBMS architecture overview
3. Layers of DBMS architecture

2. Course organization

1. Planned schedule
2. Personnel
3. Course resources
4. Exercises
5. Exam

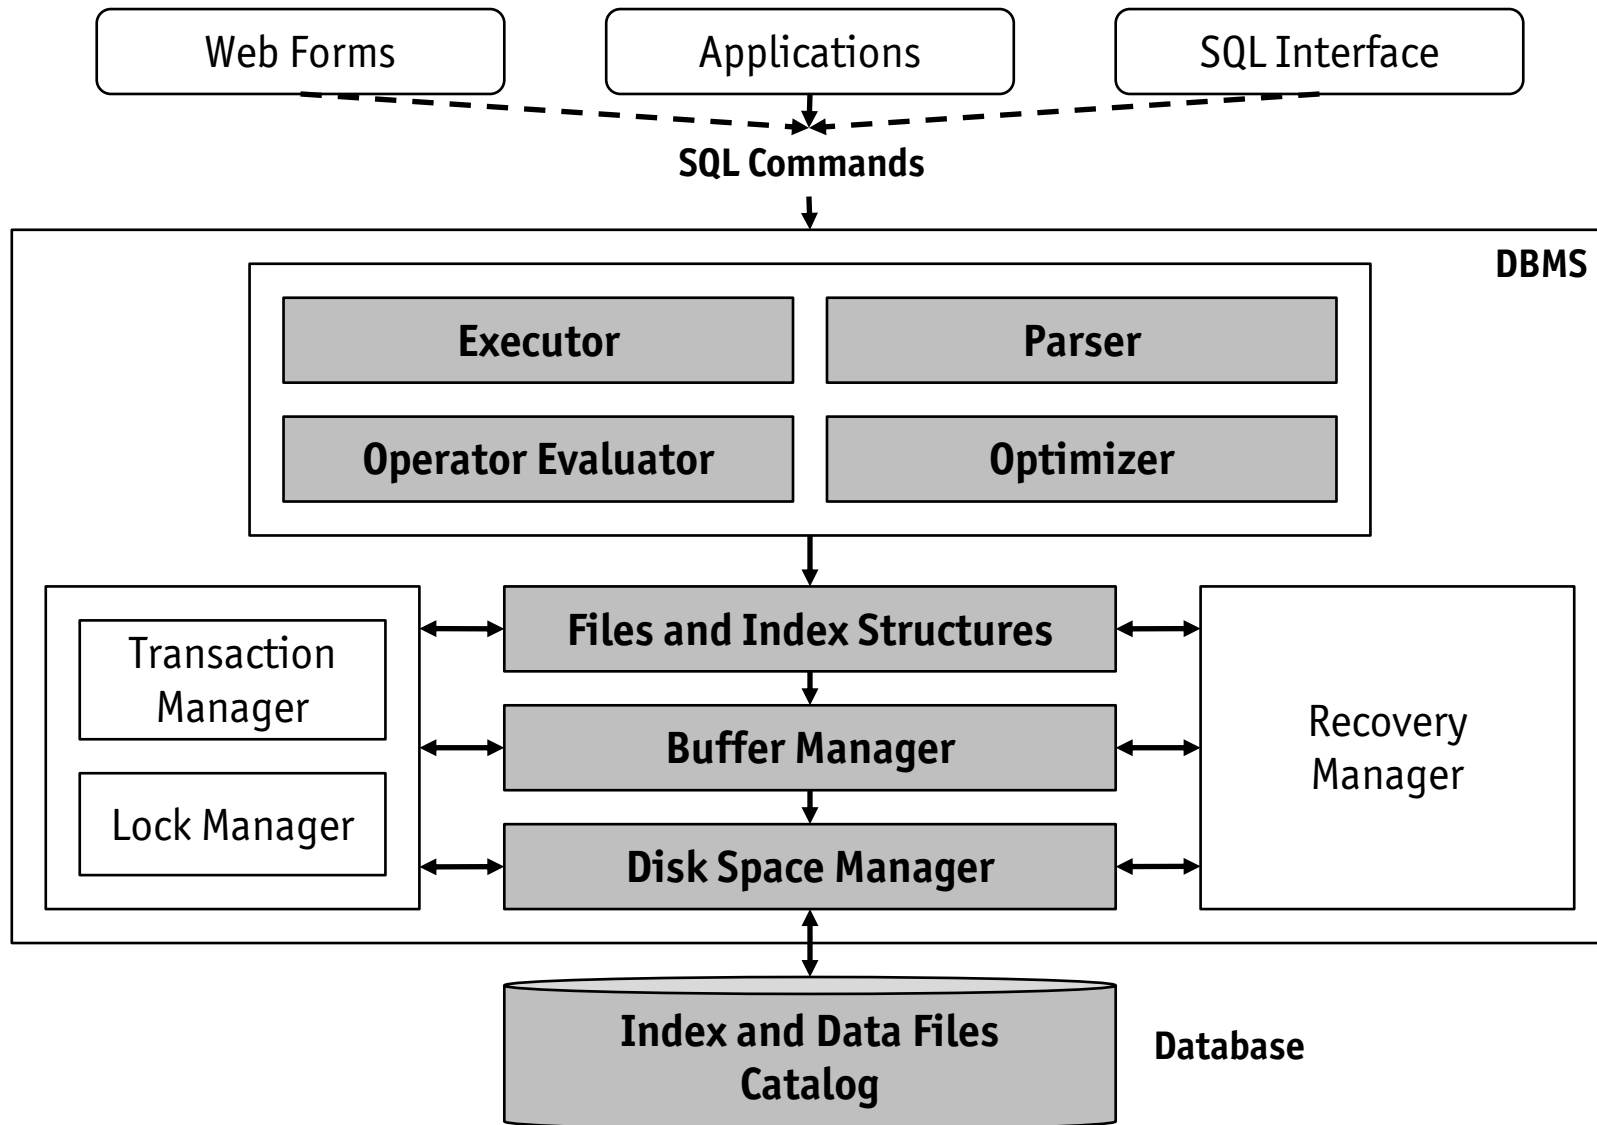
The Big Picture



Course Contents and Goals

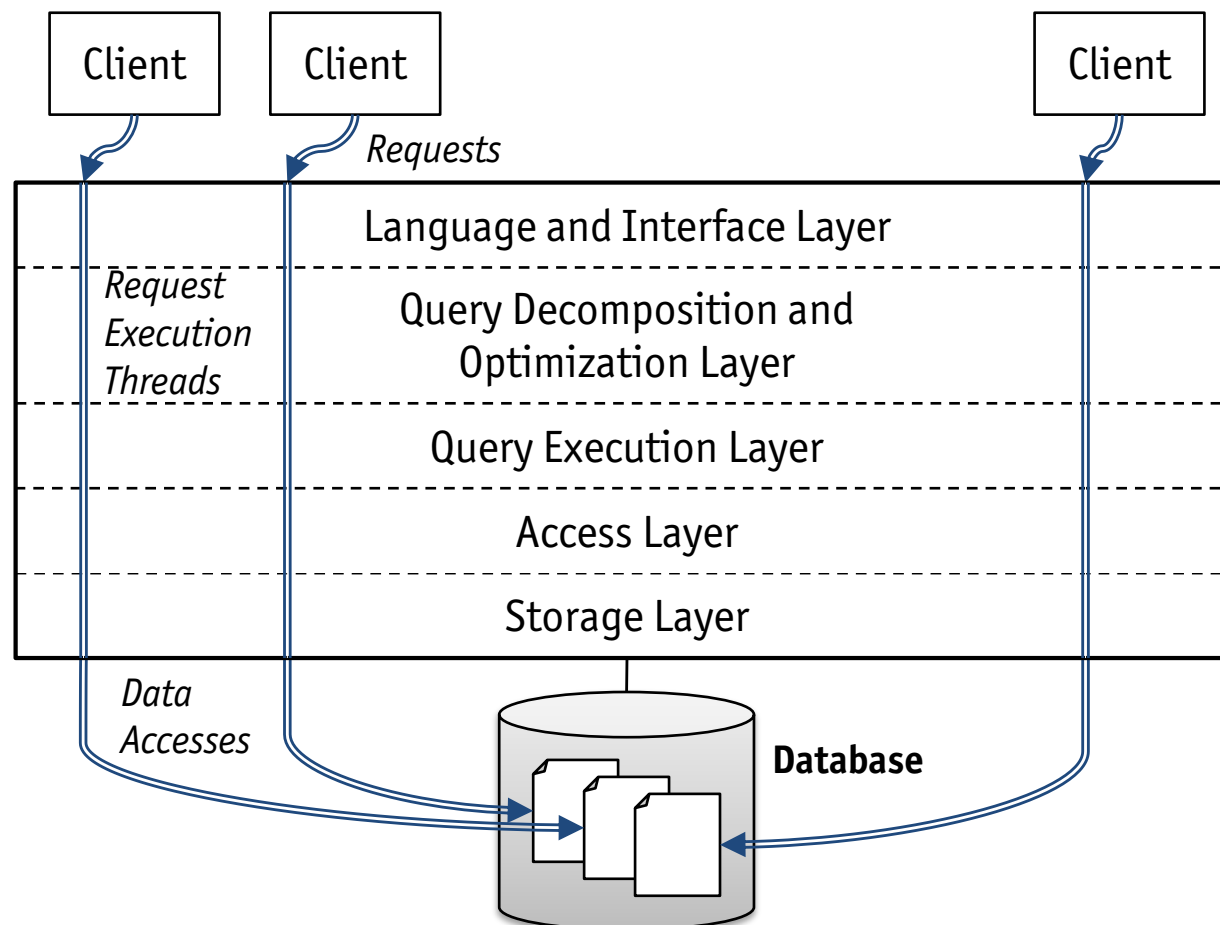
- Builds on introductory database course with focus on internals at system-level, rather than functionality at interface-level
- Learn how a DBMS can...
 - **organize and access files** on hard disks to minimize costly I/O traffic
 - **translate SQL statements** into efficient query execution plans
 - **sort/combine/filter data volumes** that exceed main memory size by far
 - be **tuned** for performance-critical applications
 - be evaluated by using **benchmarks**

DBMS Architecture



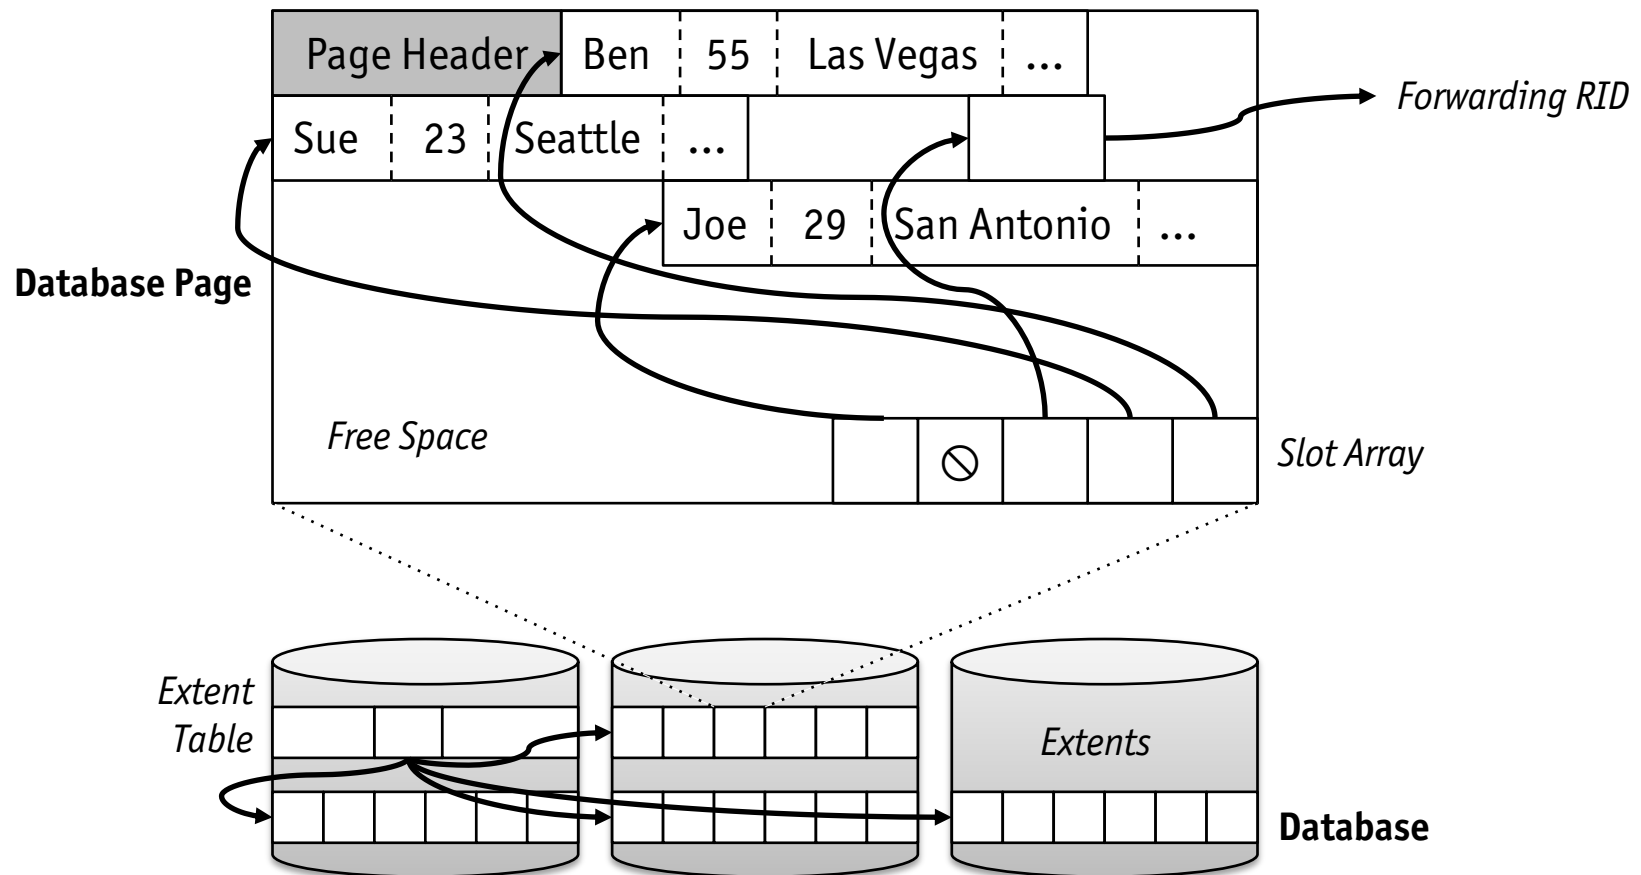
Layered DBMS Architecture

- DBMS implements its functionality in a layered architecture
 - **incrementally** add more abstractions with each layer
 - from **low-level** block I/O devices to **high-level** declarative user interface



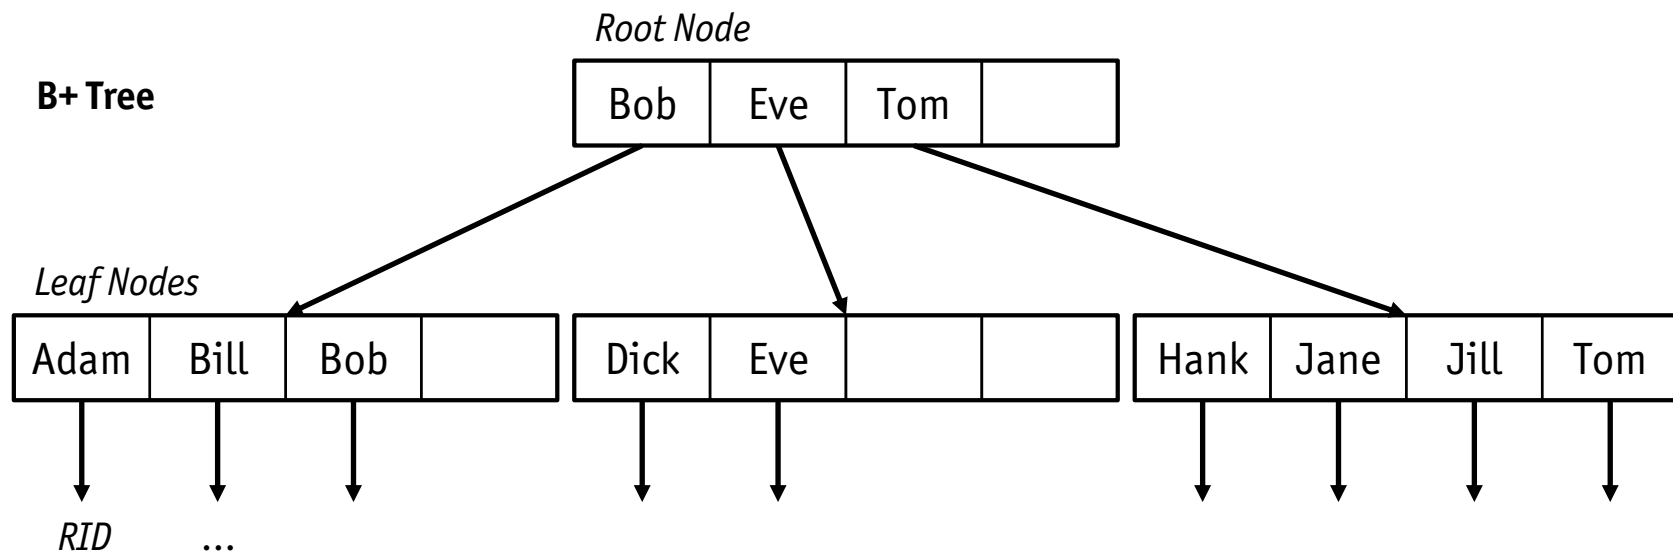
Storage Structures

- DBMS data structures are mapped to fixed-length blocks
 - basic I/O **unit of transfer** between main and secondary memory
 - true for **all** types of DBMS, i.e., relational, object-relational, etc.



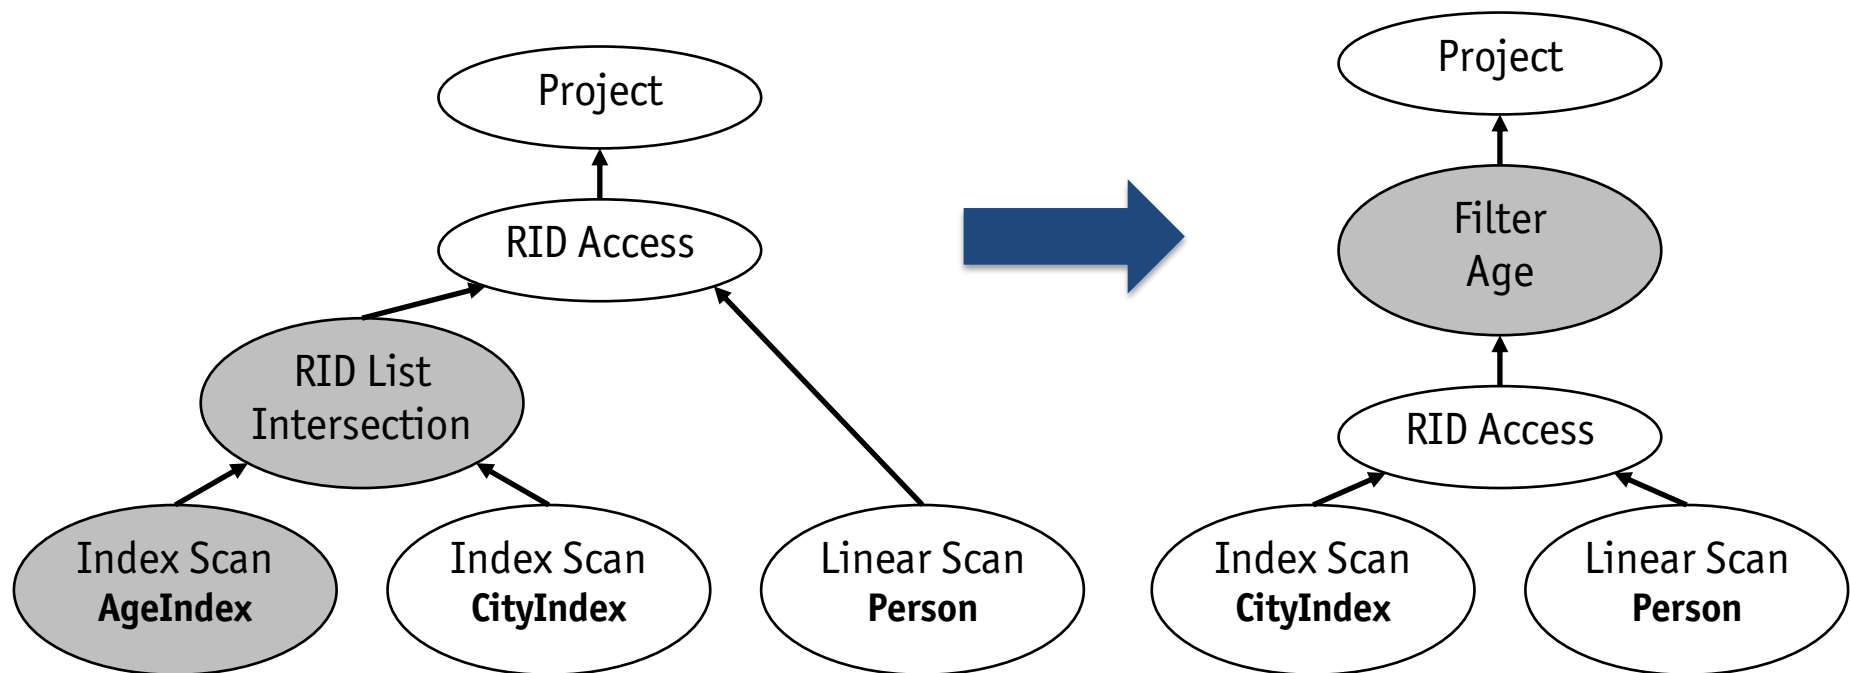
Access Paths

- DBMS indexing techniques enable fast content-based access
 - **tree-structured** and **hash-based** methods
 - **application-specific** methods



Query Processing

- DBMS processes queries and returns results
 - **declarative query specification**, e.g., expressed in SQL, is optimized and transformed into efficient **query execution plan** (QEP)
 - QEP is a **sequential** or **parallel** program that computes the query results



Planned Course Schedule

Week	Date	Topic	Exercise	Due
1	22.10.2018	Introduction	Disk Storage	30.10.2018
2	29.10.2018	Disk and Buffer Management	Disk Management	06.11.2018
3	05.11.2018	File Management	Buffer Management	20.11.2018
4	12.11.2018	Tree-structured Indexing (I)		
5	19.11.2018	Tree-structured Indexing (II)	B+ Trees	27.11.2018
6	26.11.2018	Hash-based Indexing	Disk-based Hash Indexes	11.12.2018
7	03.12.2018	Query Evaluation and Optimization		
8	10.12.2018	External Sorting	External Sorting	08.01.2019
9	17.12.2018	Evaluating Relational Operators (I)		
Christmas Break				
10	07.01.2019	Evaluating Relational Operators (II)	Relational Operators	22.01.2019
11	14.01.2019	Relational Algebra Equivalences		
12	21.01.2019	Histograms	Cost and Estimation	29.01.2019
13	28.01.2019	Nested Queries	Query Optimization	12.02.2019
14	04.02.2019	Optimizer Architectures		
15	11.02.2019	Physical Database Design and Tuning	Flashback	

Personnel

- Prof. Dr. Michael Grossniklaus

- Office PZ 806
- E-mail michael.grossniklaus@uni-konstanz.de



- Leonard Wörteler

- Office PZ 813
- E-mail leonard.woerteler@uni-konstanz.de



Curriculum Vitae

- **1996-2008:** ETH Zurich
 - Dipl. Inf.-Ing. (MSc) in Computer Science
 - Dr. sc. techn. (PhD) in Moira Norrie's group
- **2008-2010:** Politecnico di Milano
 - Post-doc in Stefano Ceri's group
- **2010-2012:** Portland State University
 - Post-doc in David Maier's group
- **2012-2013:** Vienna University of Technology
 - Post-doc in Gerti Kappel's group
- **since 2013:** University of Konstanz
 - Professor for *Databases and Information Systems*




Course Resources

- Registration
 - **ZEuS** for course-related e-mail
 - **Ilias/ecoach** for participation in the exercises
 - **StudIS** for admittance to the exam
- Literature
 - lectures slides can be downloaded from the course website
 - textbook (see next slide)
 - further readings will be published on the course website

Textbook

- *Raghu Ramakrishnan and Johannes Gehrke*
Database Management Systems (3rd Edition)
McGraw-Hill, 2002
 - course will focus on chapters 8 to 15 and 20
- Ten copies of the **2nd edition** of the book are available at the university library
 - not much difference for the scope of this course
 - course will focus on chapters 7 to 14 and 16



Prerequisites

- The following skills are **mandatory** prerequisites to attend and successfully complete this course
- **Foundations of database systems** (INF-12040 or equivalent)
 - E/R model and conceptual design, relational model, relational algebra and calculus, query languages, database application programming, etc.
- **Computer systems** (INF-11740, INF-11880, or equivalent)
 - computer architecture, operating systems, compilers, networks, etc.
- **System programming** (INF-11930 or equivalent)
 - some of the exercises are based on a (somewhat) complex Java project
 - students **must** have the ability to program in Java (read, understand, design, and write **high-quality** Java code)
- **Key competences** (INF-10175 or equivalent)
 - Subversion, LaTeX, etc.

Rooms and Dates

- Lecture

- Monday 3:15 pm – 4:45 pm Room P 602
- Tuesday 3:15 pm – 4:45 pm Room P 1138

- Exercise

- Wednesday 5:00 pm – 5:45 pm Room P 1138

Exercises

- Assignments will be made available on Ilias
 - solved in pairs of two students
- Five two-week programming projects
 - based on Minibase for Java
 - opportunity to implement core components of a DBMS yourself
- Grading
 - one-week assignments give 20 points
 - two-week assignments give 40 points
 - 50% of total course grade (if exam is passed)

Assignments and Exam Grading

- Same grading scale used for assignments and exam
- Exam will consist of four exercises with 25 points each
 1. disk and file storage
 2. hash and tree-based indexes
 3. query evaluation
 4. query optimization
- Overflow bonus points to exam points conversion

$$P_{exam} = \left\lceil \frac{1}{2} \cdot \frac{100}{280} \cdot P_{bonus} \right\rceil$$

Percentage of Points	German Grade	Letter Grade
100-91	1.0	A
90-86	1.3	A-
85-81	1.7	B+
80-76	2.0	B
75-71	2.3	B-
70-66	2.7	C+
65-61	3.0	C
60-56	3.3	C-
55-51	3.7	D+
50-45	4.0	D
44-0	5.0	F

Exam

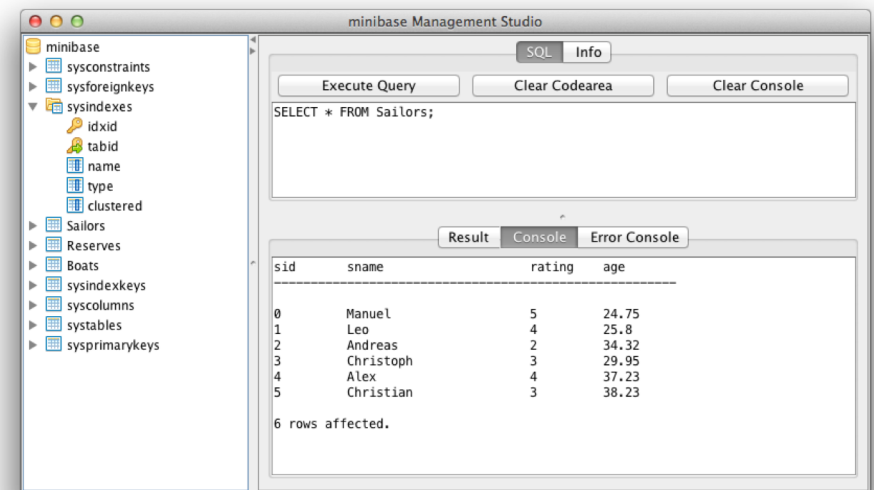
- Written exam
 - two-hour exam
 - closed-book exam, but one hand-written double-sided piece of A4 paper with notes is permitted
- Content
 - lecture and lecture slides
 - (written) exercises
 - relevant chapters from textbook
- Scheduling
 - first date February 26, 2019 (C 353)
 - second date April 9, 2019 (C 353)
- Grading
 - 50% of total course grade (if exercises are passed)

Minibase!


- Complete single-user database management system
 - originally developed in C/C++ by Ramakrishnan/Gehrke at Wisconsin
 - ported to Java at Purdue and Portland State
 - extended and improved at Konstanz

- **New features**

- new data types
- different buffer replacement policies
- additional indexes and operators
- more compliance to SQL standard
- cascades-style query optimizer
- simple user interface



- <https://github.com/DBIS-UniKN/minibase> (private, obviously)



Database System Architecture and Implementation

TO BE CONTINUED...